

APAI2024 - LAB02

PULP_Embedded_Programming

Authors: Davide Nadalini, Lorenzo Lamberti, Alberto Dequino, Manuele Rusci, Francesco Conti

Contacts: d.nadalini@unibo.it, lorenzo.lamberti@unibo.it, alberto.dequino@unibo.it

Links: [GitHub Link \(code\)](#), [GDoc Link \(solution template\)](#)

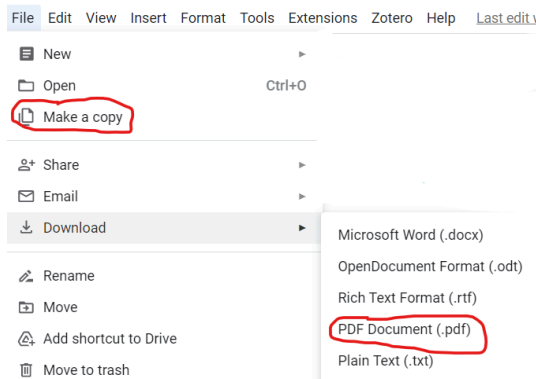
Summary

1. Subject(s):
 - PULP architecture, vector sum, matrix-vector mul, profiling code execution;
2. Programming Language: C;
3. Lab duration: 3h
4. Objective: Embedded programming & profiling
you will learn basics of embedded programming, the pulp architecture, basic operations (sum & matmul), and how to profile your code execution (MAC, cycles) !
5. Assignment:
 - 2 tasks
 - Time for delivery: 1 week
 - **Submission deadline: Oct 18th 2024**

How to deliver the assignment

You will deliver ONLY THIS TEXT FILE, no code

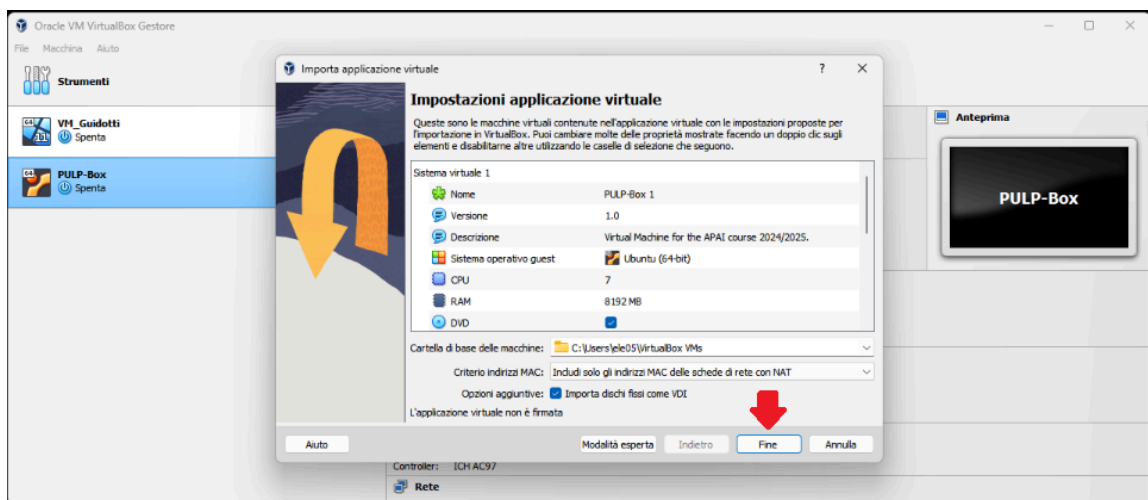
- Copy this google doc to your drive, so that you can modify it. (File -> make a copy)
- Fill the tasks on this google doc.
- Export to pdf format.
- Rename the file to: LAB<number_of_the_lesson>_APAI_<your_name>.pdf
- Use Virtuale platform to load ONLY your .pdf file



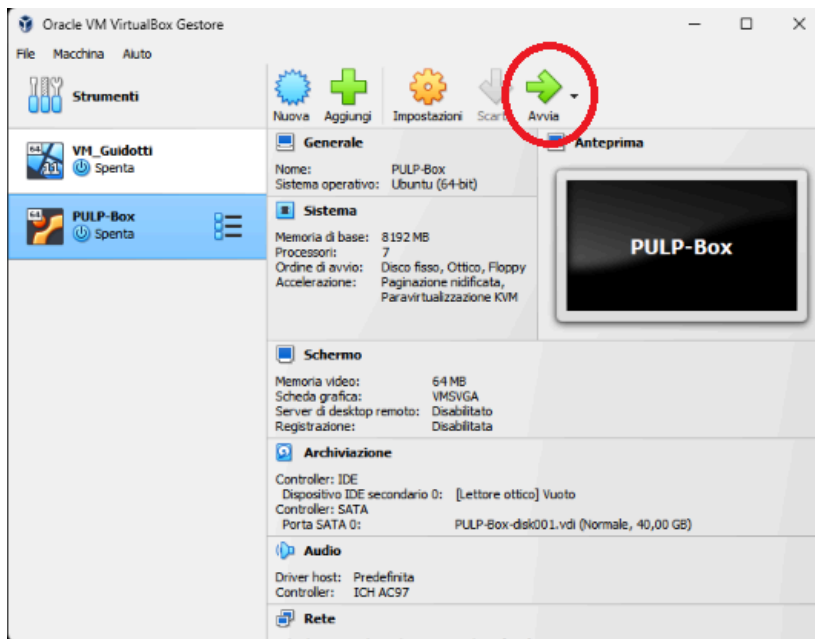
LAB STARTS HERE

0. Access to the local VM and setup pulp-sdk

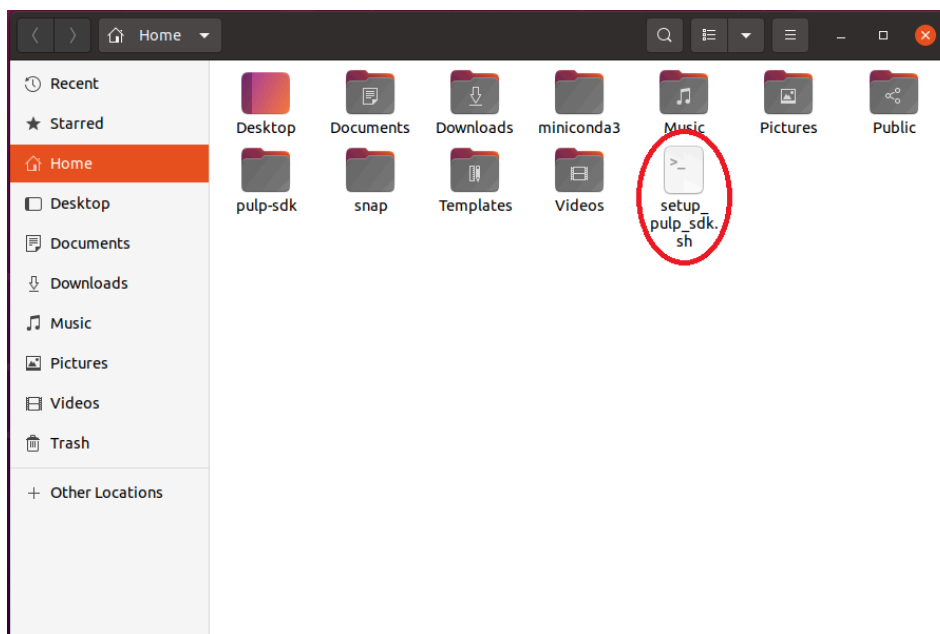
- On the lab's PCs, open the file explorer and go to This PC, C:/VM_Nadalini
- Double click on PULP-box.ova
- VirtualBox opens, just click on “Fine”



- Wait for the VM to be imported
- Open the VM with “Avvia”



- Password is 'pulp'
- Open a terminal (right click – open a new terminal)
- Setup the PULP-SDK: `source setup_pulp_sdk.sh`



- Clone GitHub repository of today's lab: `git clone https://github.com/EEESlab/APAI24-LAB02-PULP-Embedded-Programming`
- `cd APAI24-LAB02-PULP-Embedded-Programming`
- `cd pulp-helloworld`
- `make clean all run`

If the output is a “Hello World!”, then everything is correctly set up!

Task 1: vector sum

0. Setup:

- Switch to the main folder: `cd`
- Open VSCode: `code .`
- Go to “vector_sum/” folder
- Every time you want to run the code, **SAVE your file** and write in the terminal “**make clean all run**”

1.1. Define N to 50 (N= vector size) and run the code:

	Question	Answer
1	What's the result of the vector_sum() function?	
2	Is the checksum correct ?	
3	Print all the elements of “array_1”. What's the output?	

Tips:

- **Question 3:** to print the element of the array, uncomment the function call “`print_array(array_1, N)`” inside the `main()`

1.2. Define N to 350 (N= vector size) and run the code:

	Question	Answer
1	What's the result of the vector_sum() function?	
2	Is the checksum correct ?	
3	Print all the elements of "array_1". What's the output?	
4	Array_1 should be filled with increasing values from 1 to N. Why isn't the case here?	

Tips:

- **Question 3:** to print the element of the array, uncomment the function call "print_array(array_1, N)" inside the main()
- **Question4:** what range of values a "char" data type can represent?

1.3. Fix the issue: cast "array_1" and the function's arguments to int (or short int)

Write down your solution


before	After your fix
<code>unsigned char array_1[N];</code>	
<code>int init_array(unsigned char * A_ar, int size)</code>	
<code>void print_array(unsigned char * A_ar, int size)</code>	

```
int vector_sum(unsigned char * A_ar, int size)
```

After the fix, answer again

	Question	Answer
1	What's the result of the vector_sum() function?	
2	Is the checksum correct ?	
3	Print all the elements of "array_1". What's the output?	

Task 2: matrix-vector product & profiling



0. Setup:

- Go to "matrix_vector_product/" folder
- Every time you want to run the code, **SAVE your file** and write in the terminal **"make clean all run"**

2.1. Implement missing code

GOAL: We want to count how many Multiply-ACcumulate operations we perform for a GEMV

TASK: Increment a counter every time we perform a MAC inside the GEMV inner loop. I prepared a variable called `mac_counter`. This is the reference code

```
int mac_counter=0;
/* generic matrix-vector multiplication */
int __attribute__((noinline)) gemv(int size_N, int size_M, int *mat_i, int *vec_i, int *vec_o)
{
    for (int i = 0; i < size_N; i++) // outer loop of the gemv
    {
        for (int j = 0; j < size_M; j++) // inner loop of the gemv
        {
            // multiply accumulate operation (MAC)
            vec_o[i] += mat_i[i * size_M + j] * vec_i[j];
            /*(vec_o+i) += *(mat_i+i*M+j)*(vec_i+j); // try to uncomment this and comment the above line
        }
    }
}
```

Put your solution below (code)

[\[HERE\]](#)

2.2. Implement missing code part II

Add performance counters to profile the gemv. The gemv is calculated with this line of code in the `main()`:

```
gemv(N, M, matrix, vector, output_vec);
```

To profile it, exploit these two functions I prepared for you. You should start the profiling right before and stop it right after

```
start_perf_counter();
stop_perf_counter();
```

Put your solution below (code)

[\[HERE\]](#)

2.3. Implement missing code part III

Enable the performance counters of our interest. We want to profile:

- Execution cycles (total)
- N° instructions executed

Here's the full list of the performance counters

```
typedef enum {
    PI_PERF_CYCLES      = 17, /*!< Total number of cycles (also includes the
        cycles where the core is sleeping). Be careful that this event is using a
        timer shared within the cluster, so resetting, starting or stopping it on
        one core will impact other cores of the same cluster. */
    PI_PERF_ACTIVE_CYCLES = 0, /*!< Counts the number of cycles the core was
        active (not sleeping). */
    PI_PERF_INSTR        = 1, /*!< Counts the number of instructions executed.
        */
    PI_PERF_LD_STALL     = 2, /*!< Number of load data hazards. */
    PI_PERF_JR_STALL     = 3, /*!< Number of jump register data hazards. */
    PI_PERF_IMISS        = 4, /*!< Cycles waiting for instruction fetches, i.e.
        number of instructions wasted due to non-ideal caching. */
    PI_PERF_LD           = 5, /*!< Number of data memory loads executed.
        Misaligned accesses are counted twice. */
    PI_PERF_ST           = 6, /*!< Number of data memory stores executed.
        Misaligned accesses are counted twice. */
    PI_PERF_JUMP         = 7, /*!< Number of unconditional jumps (j, jal, jr,
        jalr). */
    PI_PERF_BRANCH       = 8, /*!< Number of branches. Counts both taken and
        not taken branches. */
    PI_PERF_BTAKEN       = 9, /*!< Number of taken branches. */
    PI_PERF_RVC           = 10, /*!< Number of compressed instructions
        executed. */
    PI_PERF_LD_EXT       = 12, /*!< Number of memory loads to EXT executed.
        Misaligned accesses are counted twice. Every non-TCDM access is considered
        external (cluster only). */
    PI_PERF_ST_EXT       = 13, /*!< Number of memory stores to EXT executed.
        Misaligned accesses are counted twice. Every non-TCDM access is considered
        external (cluster only). */
    PI_PERF_LD_EXT_CYC   = 14, /*!< Cycles used for memory loads to EXT.
        Every non-TCDM access is considered external (cluster only). */
    PI_PERF_ST_EXT_CYC   = 15, /*!< Cycles used for memory stores to EXT.
        Every non-TCDM access is considered external (cluster only). */
    PI_PERF_TCDM_CONT    = 16, /*!< Cycles wasted due to TCDM/log-interconnect
        contention (cluster only). */
} pi_perf_event_e;
```

Ref: `/rtos/pmsis/pmsis_api/include/pmsis/chips/default.h`

You can enable them in the `start_perf_counter()` function.

Complete the code where you find `/* YOUR CODE HERE */` with the right performance counters (see previous figure)


```

void start_perf_counter()
{
    // enable the perf counter of interest
    pi_perf_conf( 1 << PI_PERF_/* YOUR CODE HERE*/ | // count cycles
                 1 << PI_PERF_/* YOUR CODE HERE*/ ); // count instructions
    // reset the perf counters
    pi_perf_reset();
    // start the perf counter
    pi_perf_start();
}

```

Put your solution below (code)

[HERE]

2.4. Implement missing code part IV

GOAL: we now are able to measure Cycles & instructions (thanks to performance counters), and MAC operations (your implementation in the code)

TASK: calculate and print the following metrics in the code (fill in where you read /*YOUR CODE HERE*/)

- CPI = cycles/instructions
- MAC/cycles
- Instructions/Cycles
- Instructions/MACs

```

/*
TASK 2.2
Measure:
- How many MAC operations are needed for the gemv
Calculate:
- CPI
- MACs/Cycles
- Instructions/Cycles
- Instructions/ MACs
*/

// N° Multiply Accumulate Operations (MACs)
/* already done with task 2.1. use the mac_counter global variable */
// CPI = cycles / n°instructions_executed
float cpi = /* YOUR CODE HERE*/
// MAC/Cycles
float mac_on_cycles = /* YOUR CODE HERE*/
// Instructions/Cycles
float instructions_on_cycles = /* YOUR CODE HERE*/
// Instructions/MAC
float instructions_on_mac = /* YOUR CODE HERE*/

```

Put your solution below (code)

[HERE]

2. The size of the matrix is $N \times M = 50 \times 50$ and the vector is $M = 50$. Run the code and fill the table:

Note:

- Profile the matrix-vector product with different compiler optimizations: -O1, -O3, -O3 with HW Loops (default is -O1)

	-O1	-O3 -mnohwloops	-O3
Cycles			
N°Instructions			
MACs			
CPI			
Instructions/Cycles			
Instructions/MACs			